

# Memory Management

## Chapter 3: Section 3.4

# Page Replacement

- What if we are out of frames?
  - Have to discard one
  - May have pages no longer used
- How to identify unneeded?

# Working Set

- Can't assume all the pages we need will be in memory when we need them.
- **Working set** - pages a process is referencing in a short period, e.g the set of pages a process is currently using
- **Sliding window** - fixed interval over which the working set is measured.
- **Page replacement** - removing a page we may not need anymore.

# Working Set

Page reference string:

1 2 1 5 7 1 6 3 7 1 6 4 2 7

We never  
reference page 5  
again after step 4

Working set		Event Number
1	1	
12	2	
12	3	
125	4	
1257	5	
1257	6	
1567	7	
1367	8	
1367	9	
1367	10	
1367	11	
1467	12	
1246	13	
2467	14	

How about removing 5 after it falls out of the 4 step sliding window?

# Working Set

Page reference string:

1 2 1 5 7 1 6 3 7 1 6 4 2 7

2 is not referenced  
between steps 2 and  
13


Working set		Event Number
1	1	1
12	2	2
12	3	3
125	4	4
1257	5	5
1257	6	6
1567	7	7
1367	8	8
1367	9	9
1367	10	10
1367	11	11
1467	12	12
1246	13	13
2467	14	14

How about removing 5 after it falls out of the 4 step sliding window?

# Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Restart the process

# Page Replacement



# Page Replacement Algorithms

- Want lowest page-fault rate
- Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string



# Optimal Algorithm

Replace page that will not be used for longest period of time

4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1
2
3
4

4

6 page faults

5

How do you know this?

Used for measuring how well your algorithm performs against a best-case

# Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		2		2		2						7		
	0	0	0		0		4		0		0						0		
		1	1		3		3		3		1						1		

page frames

# Optimal

- Unrealizable
- Have to be able to read the future
- Provides a best-case

# Not Recently Used

- Each page has two status bit, R (read) and M (modified)
- On startup all bits are set to 0.
- Every page reference the R bit is set, every dirty page the M bit is set
- Every clock tick R bits are cleared

# Not Recently Used

- On a page fault, OS inspects all pages and puts them into 4 categories
- Class 0: not referenced, not modified
- Class 1: not referenced, modified
- Class 2: referenced, not modified
- Class 3: referenced, modified
- NRU removes a page at random from the lowest class


# Not Recently Used

- Better to removed a modified page than a clean page frequently used.
- Easy to understand
- Moderately efficient to implement
- Not optimal, but adequate.

# First In, First Out

- Eject the oldest page.
- Very low overhead
- Easy to implement
- Poor performance and erratic behavior.
  - Can end up evicting the most frequently used page
- Subject to Bélády's anomaly

# Graph of Page Faults Versus The Number of Frames





# Bélády's anomaly

- Increasing the number of page frames results in an increase in the number of page faults for a given memory access pattern.

Page Requests	3	2	1	0	3	2	4	3	2	1	0	4
Newest Page	3	2	1	0	3	2	4	4	4	1	0	0
	3	2	1	0	3	2	2	2	4	1	1	
Oldest Page	3	2	1	0	3	3	3	2	4	4		

3 page frames  
9 faults (red)

Page Requests	3	2	1	0	3	2	4	3	2	1	0	4
Newest Page	3	2	1	0	0	0	4	3	2	1	0	4
	3	2	1	1	1	0	4	3	2	1	0	
	3	2	2	2	1	0	4	3	2	1		
Oldest Page	3	3	3	2	1	0	4	3	2			

4 page frames  
10 faults (red)

# First-In-First-Out (FIFO) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (3 pages can be in memory at a time per process)

1	1	4	5	
2	2	1	3	
3	3	2	4	

9 page faults

- 4 frames

1	1	5	4	
2	2	1	5	
3	3	2		
4	4	3		

10 page faults

- Belady's Anomaly: more frames, more page faults

# FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																
	0	0	0																
		1	1																


2	2	4	4	4	0														
3	3	3	2	2	2														
1	0	0	0	3	3														

0	0																		
1	1																		
3	2																		


7	7	7																	
1	0	0																	
2	2	1																	

page frames

# FIFO Illustrating Belady's Anomaly




# Second Chance Algorithm



- Sorted in reference time order.
- When page referenced, set R bit
- When faulted check oldest page, if R bit set then set page time to current time, clear R bit and move to the end.
- Repeat until page found with no R bit set.

# Clock Page Algorithm



When a page fault occurs,  
the page the hand is  
pointing to is inspected.  
The action taken depends  
on the R bit:

R = 0: Evict the page

R = 1: Clear R and advance hand

# Least Recently Used (LRU) Algorithm

- Reference string: 1, 2, 3, 4, 1, 2, **5**, 1, 2, **3**, **4**, **5**

1	1	1	1	<b>5</b>
2	2	2	2	2
3	<b>5</b>	5	<b>4</b>	4
4	4	<b>3</b>	3	3

- Impractical to track.
  - Would need one additional memory access for each memory access or
  - Linked list of each page with most recent in front.
- Can approximate, similar to NRU
  - Need hardware assistance
  - Page reference bit in page table

# LRU Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2				2		4	4	4	0				1		1		1
	0	0	0				0		0	0	3	3				3		0		0
		1	1				3		3	2	2	2				2		2		7

page frames



# LRU Algorithm

- Stack implementation – keep a stack of page numbers in a double link form
  - Page referenced:
    - move it to the top
    - requires 6 pointers to be changed
- No search for replacement

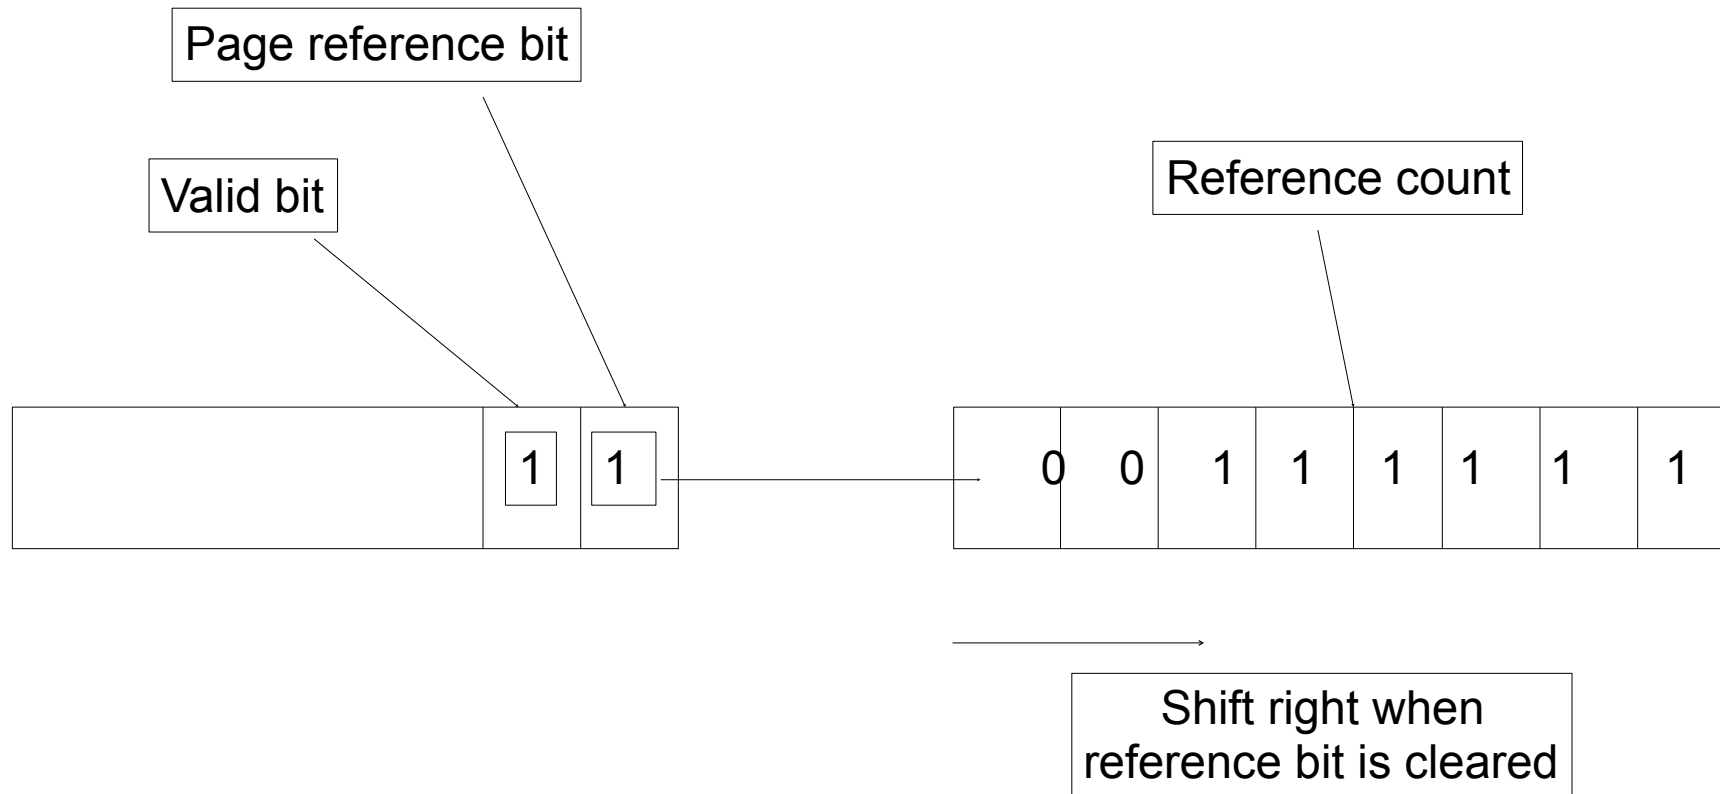
# LRU Algorithm

- Stack implementation – keep a stack of page numbers in a double link form
  - Page referenced:
    - move it to the top
    - requires 6 pointers to be changed
- No search for replacement

# Reference Counter Variation

- Also known as Aging
- As reference bits are cleared, shift into a counter
  - one for each page
- Smallest counter referenced is least recently

# Page Reference Bit & Counter



# Page Reference Bit & Counter

	R bits for pages 0-5, clock tick 0	R bits for pages 0-5, clock tick 1	R bits for pages 0-5, clock tick 2	R bits for pages 0-5, clock tick 3	R bits for pages 0-5, clock tick 4
	1 0 1 0 1 1	1 1 0 0 1 0	1 1 0 1 0 1	1 0 0 0 1 0	0 1 1 0 0 0
Page					
0	10000000	11000000	11100000	11110000	01111000
1	00000000	10000000	11000000	01100000	10110000
2	10000000	01000000	00100000	00010000	10010000
3	00000000	00000000	10000000	01000000	00100000
4	10000000	11000000	01100000	10110000	01011000
5	10000000	01000000	10100000	01010000	00101000
	(a)	(b)	(c)	(d)	(e)

# Counting Algorithms

- Keep a counter of the number of references that have been made to each page
- **LFU Algorithm**: replaces page with smallest count
- **MFU Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

# Allocation of Frames

- Each process needs *minimum* number of pages
- Example: IBM 370 – 6 pages to handle SS MOVE instruction:
  - instruction is 6 bytes, might span 2 pages
  - 2 pages to handle *from*
  - 2 pages to handle *to*
- Two major allocation schemes
  - fixed allocation
  - priority allocation

# Fixed Allocation

- Equal allocation – For example, if there are 100 frames and 5 processes, give each process 20 frames.
- Proportional allocation – Allocate according to the size of process

—  $s_i$  = size of process  $p_i$

—  $S = \sum s_i$

—  $m$  = total number of frames

—  $a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$

$$m = 64$$

$$s_1 = 10$$

$$s_2 = 127$$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$



# Priority Allocation

- Use a proportional allocation scheme using priorities rather than size
- If process  $P_i$  generates a page fault,
  - select for replacement one of its frames
  - select for replacement a frame from a process with lower priority number


# Global vs. Local Allocation

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
- **Local replacement** – each process selects from only its own set of allocated frames

# Thrashing

- If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
  - low CPU utilization
  - operating system thinks that it needs to increase the degree of multiprogramming
  - another process added to the system
- **Thrashing** – a process is busy swapping pages in and out

## Thrashing (Cont.)



# Demand Paging and Thrashing

Why does demand paging work?

Locality model


Process migrates from one locality to another

Localities may overlap

Why does thrashing occur?

$\Sigma$  size of locality > total memory size

# Locality In A Memory-Reference Pattern



# Page-Fault Frequency Scheme

- Establish “acceptable” page-fault rate
  - If actual rate too low, process loses frame
  - If actual rate too high, process gains frame

